

CSE541 Class 1

Jeremy Buhler

January 14, 2019

<https://classes.engineering.wustl.edu/cse541t/>

1 Welcome and Administritivia

Welcome to CSE 541!

- This is a course about algorithms. What does that mean?
- *Not* a course designed to passively teach specific algorithms. (That was 241/247.)
- Rather, a course for you to learn how to design and analyze your *own* algorithms.
- Course is designed to teach these skills.
- “Active learning” – no good unless you engage with the material!!!

How is this course structured?

- **Class Meetings:** explanations and illustrative examples of algorithmic techniques that work for many problems.
- **Practice Problems:** ungraded problems with solutions to reinforce understanding of techniques
- **Recitation:** led by TA, to develop proof skills in an interactive setting and improve understanding of material
- **Homeworks** (40%): graded problems without solutions to see if you can apply techniques to new problems
- **Exams** (60%): final opportunity to demonstrate that you’ve mastered the techniques

What do you need to know right now?

- The course web site is your friend. All assignments and practice problems will appear there.
- Course discussions and communication with instructor and TAs is conducted through our Piazza board.

- Homeworks are submitted electronically via Canvas. Detailed instructions on formatting and submission will be posted to the website.
- Read the course overview! It has all the grody administrative details.
- There is a **collaboration policy** for homeworks (not practice probs) – tries to balance fun of collective problem solving with promotion of individual skill building and assessment.
 - Please limit group problem solving to at most 4 people at once
 - Nothing written brought to discussion or carried away
 - “Iron Chef Rule” – one hour between discussion and writeups
 - Report (as part of your homework documents) any assistance you receive.

Who is here to help you, and how do you contact us?

- TAs: Aaron, Adrien, Alex, Zihao
- To contact myself or TAs, use Piazza.
- Our office hours will all be posted on Piazza this week.

What do I expect you to know already?

- First: how to write a basic proof of correctness and/or running time
- Given a problem statement, I *always* expect you to produce
 1. An effective method for solving the problem (“algorithm”)
 2. A proof that the algorithm correctly solves the problem
 3. A proof that the algorithm is efficient (running-time analysis)
- Second: asymptotic complexity analysis
- Third: basic algorithms and data structures: hashing, sorting and searching, binary trees (including results, but not implementation, for balanced trees), elementary graph algorithms, (some) greedy algorithms
- To warm up and check your grasp of the above skills, there is an ungraded “Homework 0” on the web site. You should be able to do all the problems found there and produce suitable formal proofs.
- You can also use Homework 0 to make sure you know how to submit homework problems through Canvas.
- If you have doubts about Homework 0 or want me to look at it, please submit it no later than January 23rd.

OK, on to the fun stuff...

2 What Are We Doing Here?

We will focus on *combinatorial optimization* problems

- We need to perform constrained optimization.
- A problem typically includes a set of *constraints* (properties that a valid solution must satisfy) and an *objective function*, or measure of goodness/badness, for solutions.
- Goal is to find a solution that is
 - *feasible* – satisfies all constraints
 - *optimal* – maximizes (or minimizes) objective
- Constraints are typically discrete rather than continuous (no such thing as an “almost feasible” solution)
- Ideally, we want an algorithm that finds an optimal feasible solution for any problem instance.
- Algorithms should be *efficient* (take worst-case time polynomial in their input size), and faster is better
- Shoot for at worst quadratic or cubic time; linear or $n \log n$ time is strongly preferred if possible.

How better to start than with an example?

3 A Scheduling Problem

- You manage a ginormous space telescope.
- Lots of astronomers want to use it to make observations.
- Each astronomer’s project p_i requires use of the telescope starting at a fixed time s_i (when their grant starts) and running for ℓ_i days.
- Only one project can use the telescope at a time.
- Your goal: justify your budget to NASA by scheduling as many projects as possible!
- More formally: given a set P of projects p_i , each occupying half-open time interval $[s_i, s_i + \ell_i)$,
- Choose a subset $\Pi \subseteq P$ of projects for which
 - No two projects’ intervals overlap (“conflict”);
 - The number of projects in Π is maximized.
- This is one of many variants of the *scheduling* or *activity selection* problem.

Example:

How should we solve this problem?

- **Suggestion 1:** repeatedly pick shortest unscheduled, non-conflicting project (i.e. one that does not conflict with any previously scheduled project).
- Does this strategy always yield an optimal solution? Prove or disprove.
- **Counterexample:**

- **Suggestion 2:** repeatedly pick non-conflicting project with earliest starting time.
- Does this always yield an optimal solution? Prove or disprove.
- **Counterexample:**

- **Suggestion 3:** first, label each project with number of *other* projects with which it conflicts. Then, repeatedly pick nonconflicting project with fewest total conflicts.
- Does this always yield an optimal solution? Prove or disprove.
- **Counterexample:**

Aaaaargh! We need a *principle* to stop the endless flailing!

4 An Approach That Works

What structure do all above solutions have in common?

- Repeatedly pick an element until no more feasible choices remain.

- Among all feasible choices, we always pick the one that minimizes or maximizes some property (project length, start time, # conflicts)
- Such algorithms are called *greedy*.
- As we've seen, greedy algorithms are frequently *not* optimal.
- Ah, but maybe we have been using the wrong property!

Let's take another wild guess...

- For each project p_i , define its *finishing time* f_i to be $s_i + \ell_i$.
- Repeatedly pick non-conflicting, unscheduled project with earliest finishing time.
- Here's a reasonably efficient implementation of this strategy in pseudocode.

```

SCHEDULE( $P$ )
  sort  $P$  in increasing order  $\{p_1 \dots p_n\}$  of finishing time  $f_i$ 
   $\Pi \leftarrow \{p_1\}$ 
   $j \leftarrow 1$ 
  for  $i$  in  $2..n$  do
    if  $s_i \geq f_j$ 
       $\Pi \leftarrow \Pi \cup \{p_i\}$ 
       $j \leftarrow i$ 
  return  $\Pi$ 

```

- Sorting the times requires $O(n \log n)$ time.
- Selection procedure takes $O(1)$ time for each i , so $O(n)$ overall.
- Hence, total complexity of SCHEDULE is $O(n \log n)$.
- But does it work????

5 Proving Correctness

Why should this greedy algorithm work where others failed? Three key observations do it for us:

1. **Greedy Choice:** For every problem instance P , there exists an optimal solution that includes first element \hat{p} picked by greedy algo.
2. **Inductive Structure:** After making greedy first choice \hat{p} for problem instance P , we are left with smaller subproblem P' , such that, if Π' is a feasible solution to P' , then $\Pi' \cup \{\hat{p}\}$ is a feasible solution to P .
(Colloquially, we say that subproblem P' has *no external constraints* restricting its feasible solutions.)
3. **Optimal Substructure:** If P' is subproblem left from P after greedy choice $\{\hat{p}\}$, and Π' is an optimal solution to P' , then $\Pi' \cup \{\hat{p}\}$ is an optimal solution to P .

Let's prove these properties for SCHEDULE's greedy choice.

- **Greedy Choice:** Let P be instance of scheduling problem, and let $\hat{p} \in P$ be first project picked by SCHEDULE. Then there exists an optimal solution to P that contains \hat{p} .
- **Pf:** we use an *exchange argument*.
- Let Π^* be *any optimal solution* to P .
- If $\hat{p} \in \Pi^*$, we are done.
- Otherwise, let Π' be solution obtained by removing soonest-ending project $p \in \Pi^*$ and adding \hat{p} .
- By construction, \hat{p} ends no later than p , so if p does not conflict with any later project of Π^* , neither does \hat{p} . Hence, Π' is feasible.

- Moreover, $|\Pi'| = |\Pi^*|$, so Π^* is optimal. QED

One down, two to go.

- **Inductive Structure:** After making greedy first choice \hat{p} for problem instance P , we are left with smaller subproblem P' , with no external constraints.
- After we select the first project \hat{p} , what is remaining subproblem?
- It's not just $P - \{\hat{p}\}$!
- Having selected \hat{p} , we cannot pick any other project that conflicts with it.
- Put another way, choosing \hat{p} imposes an *external constraint* on subproblem $P - \{\hat{p}\}$, because not every feasible solution to the subproblem can be combined with the greedy choice.
- *Simple fix:* define the subproblem to be

$$P' = P - \{\hat{p}\} - \{\text{projects that conflict with } \hat{p}\}.$$

Now any feasible solution to P' can feasibly be combined with \hat{p} , and so there is no external constraint on P' .

Two down, one to go.

- **Optimal Substructure:** If Π' is an optimal solution to subproblem P' , then $\Pi' \cup \{\hat{p}\}$ is an optimal solution to P .
- **Pf:** let Π' be as given.
- Then $\Pi = \Pi' \cup \{\hat{p}\}$ is a feasible solution to P , with size $|\Pi| = |\Pi'| + 1$.

- Now suppose Π were not optimal.
- Let Π^* be an optimal schedule containing \hat{p} . (Such a solution must exist by the Greedy Choice Property.)
- Then $\Pi^* - \{\hat{p}\}$ is a feasible schedule for P' with

$$|\Pi^* - \{\hat{p}\}| > |\Pi - \{\hat{p}\}| = |\Pi'|,$$

contradicting optimality of Π' .

- Conclude that Π must be optimal. QED

OK, that was fun. But why do these three facts constitute a proof that SCHEDULE always obtains an optimal solution?

- **Claim:** SCHEDULE's solution is optimal for every problem instance P .
- **Pf:** by induction on size of problem P .
- **Bas:** if P has size 1, greedy solution is trivially as good as optimal (it picks the one element).
- **Ind:** suppose SCHEDULE's solution is optimal for problem instances of size $< k$.
- Consider an instance P of size k .
- Let P' be subproblem obtained from P after making first greedy choice, and let \hat{p} be this choice. Observe that $|P'| < |P|$.
- By IH, SCHEDULE produces an optimal feasible solution Π' for P' . (This claim holds by inductive structure property.)
- Optimal substructure property guarantees that $\Pi' \cup \{\hat{p}\}$ is an optimal feasible solution for P .
- Hence, SCHEDULE optimally solves P of size k . QED

Key Observation: the inductive proof invokes the two structural properties as “subroutines”. The optimal substructure property in turn uses the greedy choice property in its proof. This form of argument is a “design pattern” for proving correctness of a greedy algorithm. It also serves as a guide to algorithm design: pick your greedy choice to satisfy G.C.P. while leaving behind a subproblem with optimal substructure!