

CSE 241 Class 15

Jeremy Buhler

October 21, 2015

Today: the dreaded balanced tree!

1 Why Balanced Trees?

We just did skiplists specifically to avoid balanced trees. Why am I subjecting you to this?

- Sometimes, you need worst-case $O(\log n)$
- You'll encounter these structures in practice (e.g. GNU collection types use red-black trees)
- In some situations, extra processing is worth the pain.
- We'll discuss one such example to motivate today's data structure.
- B-Trees!

2 Disk-bound Computations

Here's a model of computation that is common in, e.g., the database community.

- computer consists of a CPU/memory and a disk
- CPU is arbitrarily fast, and memory is arbitrarily fast but small
- disk reads and writes are very slow, but disk is large
- disk can read or write data only in fixed-sized chunks ("blocks" or "pages")
- algorithm's cost measured by # of disk operations, specifically # of pages read or written

Is this a sensible model of computation? Here are some rough numbers for modern high-performance computing systems:

storage	access time	size
L2 cache	2 ns	10^6 bytes
L3 cache	4 ns	10^7 bytes
DRAM	12 ns	10^{10} bytes
SSD	$300\mu s$	10^{12} bytes
disk	4-8 ms	10^{13} bytes

How much would you pay for 16 GB of DRAM today? About \$90. Compare that to two TB (2000 GB) of disk for around \$75.

Indexing structures for big databases are usually disk-bound.

3 Intro to Disk-based Trees

A B-tree is a type of tree adapted to be stored mostly on disk.

- Each node x can hold a variable number of keys
- If x has $n(x)$ keys, then it has $n(x) + 1$ children
- # of keys per node bounded by some maximum m
- $m + 1$ (max # of children per node) is called the **branching factor**

Here's a simple example of a B-tree – more details soon!

A few comments on how this fits into our model...

- Each node fits in one disk page (choose size accordingly)
- To access a node x , must do **disk-read**(x)
- To modify a node x , must do **disk-write**(x)
- Assume no fancy caching (disk pages must be reread on each access), except that root node is always cached in memory.

Disk properties determine tree's branching factor as follows.

- Branching factor determined by how many keys plus child pointers can be crammed into a disk page
- *Example:* suppose we can fit 999 keys into a disk page

- How many keys can we store in a tree of height 3?
- (That means root plus two levels, **not** the same as the book!)
- max # of nodes in tree given by

$$1 + 1000 + 1000^2 = 1001001.$$

- max # of keys is 999 times max # of nodes, or almost 1 billion.
- If root is cached, any node can be reached in only two disk accesses!

4 B-tree Definition

B-trees are not just any old disk-based tree. They have special properties designed to maintain desirable ratio of size to height.

Defn: A B-tree T is a rooted tree (with root $\text{root}[T]$) with the following properties:

1. Every node x in the tree has the following info:
 - $n(x)$ – # of keys stored in x
 - stored keys $k_1(x) \dots k_{n(x)}(x)$ in sorted order
 - $\text{leaf}(x)$ – true iff x is a leaf node
2. Every non-leaf node contains $n(x) + 1$ child pointers $c_1(x), \dots, c_{n(x)+1}(x)$
3. Every key in subtree rooted at $c_i(x)$ lies between $k_{i-1}(x)$ and $k_i(x)$. End ranges are half-open.
4. Every leaf has same depth, which is tree height h .
5. We fix a *minimum degree* $t \geq 2$ for the tree.
 - Every node may have at most $2t - 1$ keys. Max branching factor for tree is therefore $2t$.
 - Any node with exactly $2t - 1$ keys is called **full** (notice that a full node always has an odd # of keys.)
 - Every node except root must have at least $t - 1$ keys. (Root may have as little as 1 key.)

For example, if $t = 2$, every internal node has at least **how many?** [2] children and at most **how many?** [4] children. This special case is called a *2-3-4 tree*.

5 B-tree Height

Does the funky minimum degree property guarantee a good ratio of height to size? (i.e. worst-case height $O(\log n)$?)

Thm: If $n \geq 1$, then any n -key B-tree of minimum degree $t \geq 2$ has height

$$h \leq \log_t \left(\frac{n+1}{2} \right) + 1 = \frac{\Theta(\log(n+1))}{\log(t)} + 1 = O(\log n).$$

Pf: If B-tree has height h , what is *fewest* number of nodes such a tree could have?

- In worst case, root has 1 key, all other nodes have $t - 1$ keys.
- Let's account the # of keys like a recursion tree.

- Summing over all levels of the tree, we have

$$\begin{aligned} n &\geq 1 + 2(t-1) \sum_{k=1}^{h-1} t^{k-1} \\ &= 1 + 2(t-1) \sum_{j=0}^{h-2} t^j \\ &= 1 + 2(t-1) \frac{t^{h-1} - 1}{t - 1} \\ &= 2t^{h-1} - 1 \end{aligned}$$

- Conclude that

$$\begin{aligned} 2t^{h-1} - 1 &\leq n \\ t^{h-1} &\leq \frac{n+1}{2} \\ h &\leq \log_t \left(\frac{n+1}{2} \right) + 1 \\ &= O(\log n) \end{aligned}$$

which is what we want. QED