

CSE 241 Class 16

Jeremy Buhler

October 26, 2015

Today: B-Trees Part Deux

1 B-Tree Search

Finding a key in a B-tree is easy

- Start at root
- If current node contains desired key, return it.
- Otherwise, determine which subtree would have key and recur on it
- Looks at only $O(h)$ nodes

Try it on example tree (find H, S, and A)

2 B-Tree Insert and Splitting

Insertion and deletion in a B-tree are interesting because we must maintain the min- and max-degree invariants.

- What's natural **insert**(k)?
- Find leaf where k belongs and put it there
- What's wrong with simple algorithm? [**wait**]
- Leaf may already be full ($2t - 1$ keys) – adding another would violate max-degree invariant.

We can try to fix insertion by **splitting**. Splitting turns a full node into two non-full nodes.

SPLIT(x)

$k \leftarrow k_t(x)$

▷ median key

create node x_ℓ from keys $k_1(x) \dots k_{t-1}(x)$

create node x_r from keys $k_{t+1}(x) \dots k_{2t-1}(x)$

move k into parent of x

place pointers to x_ℓ and x_r to left and right of k

Example of splitting:

Can we always split a node x ?

- What if x 's parent is full?
- Would be nowhere to put median key of x ! So, let's ensure this bad case does not happen
- What if x is the root?
- Can create a new root x of size 1 to hold x 's median key
- (B-trees grow *up* from the root!)
- How do we find x 's parent? (it has no parent pointer)
- Will assume parent is cached at time of split (OK for insert, delete below)

3 Insertion Algorithm

- To avoid complications, want to visit each node on path to insertion point only once.
- Implies only one disk read per node on path, or $O(h)$ total.
- We split *preemptively* to avoid backtracking.
- Algorithm uses recursive subroutine $\text{DOINSERT}(x, k)$
- Will maintain following invariant (*):

When we call $\text{DOINSERT}(x, k)$, either x is the root, or x 's parent is not full.

```

INSERT( $T, k$ )
  DOINSERT( $\text{root}[T], k$ )

DOINSERT( $x, k$ )
  if  $x$  is full
     $m \leftarrow k_t(x)$             $\triangleright$  median key of  $x$ 
    SPLIT( $x$ )                    $\triangleright$  create  $x_\ell, x_r$ 
    if  $k < m$ 
       $x \leftarrow x_\ell$ 
    else
       $x \leftarrow x_r$ 

  if leaf( $x$ )
    place  $k$  into  $x$ 
  else
     $y \leftarrow$  correct child of  $x$ 
    DOINSERT( $y, k$ )

```

Examples? [work from the sheet]

Correctness? Argue inductively that *split never fails*. Conclude that k can be inserted at end of algorithm because we can create a non-full node if needed.

- Prove by induction on number of calls to DOINSERT.
- **Base:** Invariant (*) holds at first call to DOINSERT, since call is made on root node.
- **Ind:** Suppose invariant (*) holds after m calls; show that it will hold after $m + 1$.
- Invariant (*) holds at start of DOINSERT, so SPLIT will succeed if it happens (always room for median in parent, or new root created).
- If we call DOINSERT(y, k), y 's parent has been split if it was full, so invariant (*) is maintained.
- When we try to insert k into x , x has just been split if it was full. Hence, x is not full, and insert succeeds. QED

4 Deletion

B-tree deletion is cute but difficult to code. We'll only sketch it here.

- Two problems with removing an arbitrary key k from tree.
- First, what happens to subtrees to left and right of k ?
- Second, k 's node might have only $t - 1$ keys – removal could violate min-degree invariant.
- As for insert, will have a recursive DODELETE(x, k). Initially call DODELETE($\text{root}[T], k$).
- To keep balance, will maintain following invariant (**):

When we call $\text{DoDELETE}(x, k)$, either x is the root, or x contains at least t keys.

- Invariant (**) implies that when we do remove k from some node, it will be the root or will still have at least $t - 1$ keys after the deletion.

Assume invariant (**) is true when $\text{DoDELETE}(x, k)$ is called. Must consider three cases:

1. If x is a leaf ...

- Simply remove k from x .
- x has no children, so no subtrees to deal with.
- Invariant (**) guarantees that x has at least t keys before deleting k .

2. If x contains k (and is not a leaf) ...

- Let y and z be left and right child nodes of k in x .
- **(a)** If y has at least t keys, replace k with **pred**(k) (largest key in subtree rooted at y).

- Now, must recursively remove **pred**(k) – call $\text{DoDELETE}(y, \text{pred}(k))$
- **(b)** Else if z has at least t keys, replace k with **succ**(k) and remove **succ**(k) from subtree rooted at z .
- **(c)** Otherwise, both y and z have exactly $t - 1$ keys.
- Hence, can **unsplit** y, z to form a new node w !
- k becomes median key of w (OK to remove k from x because by invariant (**), x has at least t keys).

