

CSE 241 Class 20

Jeremy Buhler

November 11, 2015

1 Weighted Version of Shortest Paths

- BFS solves *unweighted* shortest path problem
- Every edge traversed adds one to path length
- What if edges have nonuniform weights? Let $w(u, v)$ be weight of edge (u, v)

Some intuition...

- BFS finds closest set of vertices ($d = 1$) to source, then next closest set ($d = 2$), and so on
- IOW, repeatedly process vertices closest to source.
- Tricky part was proving that every vertex is reached via a shortest path.
- Can we use the same idea in weighted case?

2 Dijkstra's Algorithm

Here is an algorithm that works when $w(u, v) \geq 0$ for all edges (u, v) .

- Uses min-first *priority queue* Q of vertices
- key is estimated distance from s to each vertex
- Initially, Q contains *all* vertices in G , but distances are unknown (∞)
- Repeatedly extract vertex x that is closest to s
- As in BFS, knowing distance from s to x tells us something about distance to x 's neighbors.
- Decrease the key of every vertex y in $\text{Adj}[x]$ to at most $d(x) + w(x, y)$.

Example

3 Pseudocode

Given graph $G = (V, E)$, starting vertex s . (Note: does not show handle manipulation)

```
DIJKSTRA( $G, s$ )  
  for  $u \in V$  do ▷ initialize  
     $u.\text{distance} \leftarrow \infty$   
     $u.\text{parent} \leftarrow \text{null}$   
     $Q.\text{insert}(u, \infty)$   
  
   $s.\text{distance} \leftarrow 0$   
   $Q.\text{decreaseKey}(s, 0)$   
  
  while  $Q$  is not empty do  
     $u \leftarrow Q.\text{extractMin}()$   
    if  $u.\text{distance} = \infty$   
      stop ▷ cannot reach any more vertices from  $s$   
    for  $v \in \text{Adj}[u]$  do  
      if  $Q.\text{decreaseKey}(v, u.\text{distance} + w(u, v))$   
         $v.\text{distance} \leftarrow u.\text{distance} + w(u, v)$   
         $v.\text{parent} \leftarrow u$ 
```

4 Running Time

- cost dominated by priority queue ops (queue size n)
- initialization: one insert per vertex (n)
- outer loop: one extractMin per vertex (n)
- inner loop: one decreaseKey per edge out of each u (m)

Hence, can write

$$T(m, n) = nT_{\text{insert}}(n) + nT_{\text{extractMin}}(n) + mT_{\text{decreaseKey}}(n)$$

- Cost depends on priority queue implementation!
- For binary heaps, all queue ops are $O(\log n)$, so

$$T(m, n) = (2n + m)O(\log n) = O(m \log n)$$

- For a Fibonacci heap, insert and decreaseKey are amortized $O(1)$
- Hence, revised run time would be

$$T(m, n) = nO(1) + nO(\log n) + mO(1) = O(n \log n + m)$$

- Is this an improvement? Yes, if graph is dense.

5 Correctness

As before, we need to show that every vertex receives its correct shortest-path distance from s . Note that u .distance never changes after u is removed from the priority queue.

Theorem: when vertex u is removed from the queue, u .distance is length of a shortest path from s to u .

- Proceed by induction on order of removal from queue.
 - **Bas:** s is removed first from queue, and it has correct distance 0.
 - **Ind:** Assume that vertex u is next to be dequeued, but it does not have its shortest-path distance.
 - Consider a shortest path p connecting s to u .
-
- s has been dequeued and u has not, so there is some last vertex x on this path that *has* already been dequeued.
 - By IH, x has its correct shortest-path distance.
 - Let y be x 's successor on path p (which has not been dequeued yet), and let p' be the prefix of p connecting s to y .
 - **Prefix p' is shortest path from s to y .** Otherwise, could replace it with a shorter path p'' , which would give a shorter path than p from s to u .
 - Hence, y received its correct shortest-path distance when x was processed, since edge $x \rightarrow y$ was explored.
 - To finish up, two possibilities:
 1. If $y = u$, then u has its correct shortest-path distance, which contradicts our assumption that this distance is wrong.
 2. If y precedes u , then y 's shortest-path distance is $\leq u$'s shortest-path distance. Hence, y 's s-p distance is strictly less than u 's current (non-s-p) distance. Conclude that y will be dequeued before u , which contradicts our assumption that u is next vertex to be dequeued.
 - Conclude that u must have its correct shortest-path distance. QED

6 Other Ways to Get Shortest Paths

Remember, Dijkstra's algorithm has an important limitation!

- Requires that $w(u, v) \geq 0$ for all edges (u, v)
- **Problem:** assumes that no prefix of a path p can have length $> p$.
- If edge weights can be negative, this assumption is violated.
- Hence, can end up dequeuing a vertex before path of least *total* weight is found.

- How could this happen? “Shortest” path could be measured in terms other than distance.
- For example, suppose that on each edge (u, v) you may be charged a fee ($w(u, v) > 0$) or paid a bonus ($w(u, v) < 0$). Goal is to find path with smallest total cost!
- In this case, you want an algorithm that deals with negative-weight edges.
- **Bellman-Ford** algorithm can do it.
- Also can detect cycles of negative weight (causes paths with arbitrarily low weight, so no “shortest”).
- Cost is $O(mn)$, which is worse than Dijkstra in general.
- *Special case:* if graph is a DAG, can reduce cost to $\Theta(m + n)$.
- Finally, suppose you need to know the shortest paths from ALL vertices to ALL vertices in G .
- If
 - you can negative-weight edges;
 - you *cannot* have negative-weight cycles (use Bellman-Ford on an augmented version of G to check!)

there is a $\Theta(n^3)$ algorithm for this problem due to Floyd and Warshall. Unless your graph is sparse, this is asymptotically faster than running Bellman-Ford once per starting vertex.

- Another algorithm for the same problem, due to Johnson, takes time $O(n^2 \log n + nm)$ when implemented with a Fibonacci heap – same as Floyd-Warshall for dense graphs, but faster for non-dense graphs.